

SHREYAS S K

GENERATIVE AI SPECIALIST

skshreyas714@gmail.com | +917019667673 | [LinkedIn](#) | [Substack](#) | [GitHub](#) | [Youtube](#) | [Portfolio](#)

SUMMARY

6+ years of experience in designing and delivering production-grade LLM systems, RAG pipelines, and agentic workflows across enterprise environments. Strong background in NLP, Computer Vision, Reinforcement Learning and ML infrastructure. Proven track record leading cross-functional initiatives, owning end-to-end delivery from architecture and implementation to documentation

CORE SKILLS

Programming : Python, PyTorch, Numpy, Pandas, Langchain, Langgraph
Technologies : Machine learning, NLP, Computer Vision, Reinforcement Learning
Cloud Expertise : AWS, GCP : ECS, EC2, S3, Sagemaker, Cloud build, Cloud Run
Generative AI : OpenAI, Gemini, Llama, Phi3, DeepSeek, Gemma, LoRA, QLoRA, RLHF, DPO
MLOps/LLMOps: Dagster, Argilla, DeepEval, Opik, Docker, Kubernetes, WandB, DeepSpeed, vLLM, accelerate, Nvidia Triton, TensorRT, Qdrant, Pixeltale

WORK EXPERIENCE

Ford Motors Pvt Limited, Bangalore – Generative AI Specialist

AI DevOps Agent – Product Owner/ Tech Anchor

Jul 2024 – Present

- Led a team of 5 to build an AI-powered DevSecOps automation platform integrating Terraform, CI/CD and security/compliance (FOSSA, 42Crunch, SonarQube, Cyclope), onboarding 150+ internal customers.
- Designed a Supervisor Agentic workflow enabling natural-language to complete deployment and compliance via MCP, Browser agent, Multimodal RAG instrumented tracing/monitoring with Opik, LLM evaluation with ConfidentialAI
- Achieved 30–60% reduction in engineering effort by automating pipelines and applying Ford-specific compliance guardrails.
- Drove product GTM internally: docs, forum demos, onboarding sessions

In-vehicle Assistant – Individual Contributor

- Developed a multi-agent system to efficiently handle vehicle controls, navigation, entertainment, vehicle info, information retrieval
- Evaluated 15+ LLMs on function calling for routing accuracy, precision/recall/F1 score, and output quality based on Hallucination for each agent
- Finetuned Llama 3.2 and Qwen 2.5 3B models and achieved ~20% improvement
- Deployed on Nvidia H100 with Triton and TensorRT optimisations to achieve <50 ms latency

Khoros India R&D Pvt Ltd, Bangalore – Machine Learning Engineer

Helix (RAG-powered Khoros Bot) – Individual Contributor

Jan 2022 – July 2024

- Built analytics engine using Taxonomy/Ontology and Topic Modeling orchestrated with Dagster; extracted entities as metadata for fine-grained answers.
- Fine-tuned retriever for domain terminology and abbreviations; fine-tuned reranker for answer relevancy.
- Reduced agent handoff rate by 38% through improved retrieval precision and content routing.

Agent Auto Assist – Individual Contributor

- Fine-tuned Mistral 7B on 1.5M call-center conversations to power smart compose; collected accept/reject feedback and applied DPO for preference alignment.
- Reduced agent response time by 61% while maintaining brand tone and compliance.

Meta, Hyderabad – ML Engineer (HCL Client Project)

Apr 2020 – Jan 2022

Test Automation CLI: Developed for Facebook's AR/VR/VA (Oculus, Ray-Ban Stories)

- Implemented image binarization with SciPy/Numba to cut OCR compute by **97%**.
- Deployed Detectron2 model for AR effects detection in WhatsApp/Messenger video calls; owned data pipeline, training, and deployment.
- Delivered DistilBERT-based textual similarity with high accuracy and low latency; CLI reduced execution time by **50%**.

HCL Technologies, Bangalore – Lead Engineer

Jul 2019 – Jan 2022

- COVID-19 X-ray classification: modified DenseNet-121 achieved 96% accuracy; quantized for edge with 50–75% size reduction and 5–15% accuracy loss.
- Android app to benchmark ARM Mali GPU inference; improved GPU inference speed by 30% (FP32) and 80% (FP16) via TFLite optimization.

- Shreyas and Kameshwar (2021). **“Diagnostic Decision Support for Medical Imaging and COVID-19 Image Classification on ARM Mali GPU”**. IEEE Global Communications Conference 2021 – Edge-AI-IoT – Published on 7th December 2021 ([link](#))
- Shreyas and Dey, (2019). **“Application of soft computing techniques in Tunneling and Underground excavations: State of the art and future prospects”**. Innovative Infrastructure Solutions. Springer Publications. Published on 25th September 2019 ([link](#))
- Shreyas and Dey, (2019). **“Comparative Efficacies of Ensemble methods and Hybrid Intelligent model in Predicting Penetration Rate of Tunnel Boring Machine”**. International IACMAG Symposium, IIT Gandhinagar 2019. Published on 7th March 2019

Indian Institute of Technology, Guwahati

Aug 2017 – June 2019

M.Tech in Civil Engineering

- Final CGPA: 9.79 / 10
- Secured 1st place in the merit list

P.E.S University, Bangalore

July 2013 – May 2017

B.Tech in Civil Engineering

- Final CGPA: 9.14 / 10
- Secured 7th place in the merit list

- **Selected in Finals** (Out of 470 teams): **Bengaluru Mobility Hackathon 2024** (Organized by IISc, IUDX, BTP) – Vehicle Re-identification across multiple cameras for O-D Flows.
- **3rd Place** (out of 1200 teams): **SEBI Hackathon** (Global Fintech Fest 2023) – Developed an AI model to detect misleading claims by financial influencers using Generative AI.
- **Top 25%** (out of 400 teams): **Autonomous Agents Hackathon** (Lablab.ai, Weaviate)
- **4th Place** (out of 60 teams): **Generative AI Hackathon** (WhatFix, AWS) – Contributed to an AI solution for social good using Generative AI.
- **Special Mentions at Vista Hackathon** for demonstrating Content Genie – An all-in-one tool for social media post creation

PUBLICATIONS

EDUCATION

ACHIEVEMENTS